

Reg

December 5, 2025

```
[29]: # =====
# CHUNK 0: MASTER SETUP & IMPORTS
# =====
import pandas as pd
import numpy as np
import time
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from statsmodels.stats.outliers_influence import variance_inflation_factor

# CONFIGURATION
# Set to None to use all data (Production), or a small number (e.g., 50000) for
# fast testing
SAMPLE_SIZE = None

print("--- 1. Loading Data ---")
df_flights = pd.read_csv('US_flights_2023.csv', low_memory=False)
df_airports = pd.read_csv('airports_geolocation.csv')
df_weather = pd.read_csv('weather_meteo_by_airport.csv')

print("--- 2. Merging & Prepping ---")
# Convert dates
df_flights['FlightDate'] = pd.to_datetime(df_flights['FlightDate'])
df_weather['time'] = pd.to_datetime(df_weather['time'])

# SAMPLING (Crucial for speed if dataset is huge)
if SAMPLE_SIZE and len(df_flights) > SAMPLE_SIZE:
    print(f"    -> Sampling {SAMPLE_SIZE} rows...")
    df_flights = df_flights.sample(n=SAMPLE_SIZE, random_state=42)
```

```

# MERGE AIRPORTS (Get Origin_State)
df_merged = df_flights.merge(
    df_airports[['IATA_CODE', 'STATE']],
    left_on='Dep_Airport',
    right_on='IATA_CODE',
    how='left'
)
df_merged = df_merged.rename(columns={'STATE': 'Origin_State'}).
    drop(columns=['IATA_CODE'])

# MERGE WEATHER (Get Rain/Wind)
df_final = df_merged.merge(
    df_weather[['airport_id', 'time', 'prcp', 'snow', 'wspd', 'tavg']],
    left_on=['Dep_Airport', 'FlightDate'],
    right_on=['airport_id', 'time'],
    how='left'
)
# Cleanup
df_final = df_final.drop(columns=['airport_id', 'time'])

print(f"SETUP COMPLETE. 'df_final' is ready. Shape: {df_final.shape}")

```

--- 1. Loading Data ---

--- 2. Merging & Prepping ---

SETUP COMPLETE. 'df_final' is ready. Shape: (6743404, 29)

[30]: # =====

```

# CHUNK 1: Feature Selection & Split
# =====
print("--- Setting up Final Regression Data ---")

# 1. FINAL FEATURES
features_final = [
    'Dep_Delay',           # Primary driver
    'Flight_Duration',     # Proxy for Distance
    'Day_Of_Week',
    'DepTime_label',       # Morning/Evening
    'Airlne',
    'Aircraft_age',
    'Origin_State',
    'tavg', 'prcp', 'wspd' # Weather factors
]
target = 'Arr_Delay'

# 2. FILTER & CLEAN
df_reg = df_final[features_final + [target]].copy()

```

```

df_reg = df_reg.dropna()

# 3. SAMPLING FOR MODELING
SAMPLE_N = SAMPLE_SIZE
if SAMPLE_N and len(df_reg) > SAMPLE_N:
    df_reg = df_reg.sample(SAMPLE_N, random_state=42)

# 4. SPLIT
X = df_reg.drop(columns=[target])
y = df_reg[target]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)

print(f"Data Ready. Training on {len(X_train)} rows, Testing on {len(X_test)} rows.")

```

--- Setting up Final Regression Data ---
Data Ready. Training on 5394723 rows, Testing on 1348681 rows.

```

[31]: # =====
# CHUNK 1.5: Phase I - Cleaning & EDA Audit
# =====
print("--- Phase I: Data Cleaning & EDA Audit ---")

# 1. DUPLICATE CHECK
initial_len = len(df_reg)
df_reg = df_reg.drop_duplicates()
print(f"Duplicates Removed: {initial_len - len(df_reg)}")

# 2. MISSING DATA CHECK
print(f"Missing Values Status: {df_reg.isna().sum().sum()} (Cleaned via
DropNA)")

# 3. OUTLIER ANALYSIS (Physics Check)
impossible_flights = df_reg[df_reg['Flight_Duration'] <= 0]
if len(impossible_flights) > 0:
    df_reg = df_reg[df_reg['Flight_Duration'] > 0]
    print(f" -> Removed {len(impossible_flights)} impossible flights (Duration
<= 0).")

# Visualization (Boxplot)
plt.figure(figsize=(10, 4))
sns.boxplot(x=df_reg['Dep_Delay'], color='orange')
plt.title('Outlier Analysis: Departure Delays')
plt.xlabel('Minutes')
plt.show()

```

```

# 4. COLLINEARITY CHECK (Correlation Matrix)
plt.figure(figsize=(10, 8))
numeric_df = df_reg.select_dtypes(include=['float64', 'int64'])
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Pearson Correlation Matrix')
plt.show()

# 5. VIF ANALYSIS
print("\n--- Syllabus Check: VIF Analysis ---")
X_vif = numeric_df.drop(columns=['Arr_Delay'], errors='ignore')
X_vif['intercept'] = 1

vif_data = pd.DataFrame()
vif_data["feature"] = X_vif.columns
vif_data["VIF"] = [variance_inflation_factor(X_vif.values, i) for i in
                   range(len(X_vif.columns))]
print(vif_data.sort_values('VIF', ascending=False))

```


--- Phase I: Data Cleaning & EDA Audit ---


Duplicates Removed: 492

Missing Values Status: 0 (Cleaned via DropNA)

-> Removed 1 impossible flights (Duration <= 0).

Outlier Analysis: Departure Delays





--- Syllabus Check: VIF Analysis ---

	feature	VIF
7	intercept	21.607212
1	Flight_Duration	1.027907
3	Aircraft_age	1.024293
6	wspd	1.022501
5	prcp	1.015051
4	tavg	1.012356
0	Dep_Delay	1.008057
2	Day_Of_Week	1.002016

[32]: # ======
CHUNK 1.6: Detailed Outlier Audit Function
======
def outlier_audit(df, features):

```

num_feats = [c for c in features if pd.api.types.is_numeric_dtype(df[c])]
cat_feats = [c for c in features if c not in num_feats]

report_rows = []
for c in num_feats:
    col = df[c]
    q1, q3 = col.quantile(0.25), col.quantile(0.75)
    iqr = q3 - q1
    lb, ub = q1 - 1.5 * iqr, q3 + 1.5 * iqr
    z = (col - col.mean()) / (col.std(ddof=0) if col.std(ddof=0) != 0 else 1)

    report_rows.append({
        "feature": c, "dtype": str(col.dtype), "count": len(col),
        "min": float(col.min()), "mean": float(col.mean()), "max": float(col.max()),
        "outliers_1.5IQR": int(((col < lb) | (col > ub)).sum()),
        "outliers_z>4": int((z.abs() > 4).sum())
    })

display(pd.DataFrame(report_rows).set_index('feature'))

# Flags for specific domain rules
flags = pd.Series(False, index=df.index)
if 'Flight_Duration' in df.columns: flags |= (df['Flight_Duration'] <= 0)
if 'Aircraft_age' in df.columns: flags |= (df['Aircraft_age'] > 80)
if 'prcp' in df.columns: flags |= (df['prcp'] > 500)
if 'wspd' in df.columns: flags |= (df['wspd'] > 200)
if 'tavg' in df.columns: flags |= (df['tavg'] < -80) | (df['tavg'] > 60)

return {"flag_index": df.index[flags]}

# Run the audit
audit = outlier_audit(X, features_final)

# Clean based on audit
X_clean = X.drop(index=audit['flag_index'])
y_clean = y.drop(index=audit['flag_index'])
print(f"Audit Complete. Rows removed: {len(audit['flag_index'])}")

```

feature		dtype	count	min	mean	max	outliers_1.5IQR	\
Dep_Delay		int64	6743404	-99.0	12.200987	4413.0	863168	
Flight_Duration		int64	6743404	0.0	140.297779	795.0	340654	
Day_Of_Week		int64	6743404	1.0	3.982793	7.0	0	
Aircraft_age		int64	6743404	1.0	13.480626	57.0	4248	
tavg		float64	6743404	-39.1	16.810826	42.2	24127	

```

prcp          float64  6743404    0.0    2.437808    571.5      1413666
wspd          float64  6743404    0.0   12.476340     59.4      146967

outliers_z>4
feature
Dep_Delay      52533
Flight_Duration 11680
Day_Of_Week      0
Aircraft_age    4248
tavg            2368
prcp            80973
wspd            11083

Audit Complete. Rows removed: 134

```

```

[33]: # =====
# CHUNK 2: The Pipeline Setup
# =====
# Identify columns
num_cols = X.select_dtypes(include=np.number).columns.tolist()
cat_cols = X.select_dtypes(exclude=np.number).columns.tolist()

# Transformer
preprocessor = ColumnTransformer(
    transformers=[
        ('num', 'passthrough', num_cols),
        ('cat', OneHotEncoder(handle_unknown='ignore', sparse_output=False), ↴
         cat_cols)
    ]
)

print("Pipeline defined.")
print(f"Numeric: {num_cols}")
print(f"Categorical: {cat_cols}")

# Quick Check on Delay Balance
delay_threshold = 15
print(f"Train Delayed > 15min: {(y_train > delay_threshold).sum()}")
print(f"Test Delayed > 15min: {(y_test > delay_threshold).sum()}")

```

```

Pipeline defined.
Numeric: ['Dep_Delay', 'Flight_Duration', 'Day_Of_Week', 'Aircraft_age', 'tavg',
'prcp', 'wspd']
Categorical: ['DepTime_label', 'Airline', 'Origin_State']
Train Delayed > 15min: 1072119
Test Delayed > 15min: 267558

```

```
[34]: # =====
# CHUNK 3: Model Training (Machine Learning)
# =====

def train_and_score(name, model):
    print(f"\nTraining {name}...")
    start = time.time()
    pipe = Pipeline(steps=[('preprocessor', preprocessor), ('model', model)])
    pipe.fit(X_train, y_train)
    preds = pipe.predict(X_test)

    rmse = np.sqrt(mean_squared_error(y_test, preds))
    r2 = r2_score(y_test, preds)
    print(f" -> RMSE: {rmse:.4f} min")
    print(f" -> R2: {r2:.4f}")
    return {'RMSE': rmse, 'R2': r2, 'Model': pipe}

# Train Baseline
results = {}
results['Linear'] = train_and_score("Linear Regression", LinearRegression())
```

Training Linear Regression...

-> RMSE: 13.5155 min
-> R2: 0.9436

```
[35]: # =====
# CHUNK 4: Diagnostics (Visuals)
# =====

final_model_pipe = results['Linear']['Model']
final_model_obj = final_model_pipe.named_steps['model']

y_pred_final = final_model_pipe.predict(X_test)
residuals = y_test - y_pred_final

plt.figure(figsize=(14, 6))

# Plot 1: Actual vs Predicted
plt.subplot(1, 2, 1)
sns.scatterplot(x=y_test, y=y_pred_final, alpha=0.3, color="#3498db")
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.title('Actual vs. Predicted')
plt.xlabel('Actual Delay'); plt.ylabel('Predicted Delay')

# Plot 2: Residuals
plt.subplot(1, 2, 2)
sns.histplot(residuals, kde=True, color="#e74c3c", bins=50)
```


```

plt.title('Residuals Distribution')
plt.xlabel('Error (min)'); plt.xlim(-50, 50)

plt.tight_layout()
plt.show()

print(f"Model R2: {r2_score(y_test, y_pred_final):.4f}")

```



Model R2: 0.9436

```

[36]: # =====
# CHUNK 5: Phase II - Statistical Inference & Stepwise Regression
# =====
print("--- Phase II: Statistical Analysis (Syllabus Requirements) ---")

# 1. PREP DATA
STATS_SAMPLE = 100000
df_stats = df_reg.sample(n=min(len(df_reg), STATS_SAMPLE), random_state=42).
    ↪copy()
df_stats_encoded = pd.get_dummies(df_stats, columns=cat_cols, drop_first=True,
    ↪dtype=int)

y_sm = df_stats_encoded['Arr_Delay'].astype(float)
X_sm = df_stats_encoded.drop(columns=['Arr_Delay']).astype(float)
X_sm = sm.add_constant(X_sm)

# 2. STEPWISE REGRESSION (Backward Elimination)
def backward_elimination(X, y, sl=0.05):
    features = X.columns.tolist()
    while len(features) > 0:
        X_curr = X[features]

```

```

    if X_curr.empty: break
    model = sm.OLS(y, X_curr).fit()
    p_values = model.pvalues
    max_p = p_values.max()
    if max_p > sl:
        worst = p_values.idxmax()
        if worst == 'const' and len(features) > 1:
            # Try to find next worst if const is worst
            worst = p_values.drop('const').idxmax()
            if p_values[worst] <= sl: break
        features.remove(worst)
    else:
        break
return model, features

print(f"Running Stepwise Regression on {len(X_sm.columns)} features...")
final_sm_model, selected_features = backward_elimination(X_sm, y_sm)

# 3. OUTPUTS
print(final_sm_model.summary())

# 4. METRICS TABLE
metrics_df = pd.DataFrame({
    "Metric": ["Adjusted R-Squared", "AIC", "BIC", "MSE"],
    "Value": [final_sm_model.rsquared_adj, final_sm_model.aic, final_sm_model.
              bic, final_sm_model.mse_resid]
})
display(metrics_df)

# 5. REQUIRED PLOT: Train vs Test vs Predicted
plt.figure(figsize=(10, 6))
plot_n = 50
viz_df = df_stats.sample(n=plot_n*2, random_state=99).copy()
viz_encoded = pd.get_dummies(viz_df, columns=cat_cols, drop_first=True, u
                             dtype=int)
viz_encoded = sm.add_constant(viz_encoded).astype(float)
viz_encoded = viz_encoded.reindex(columns=selected_features, fill_value=0)

viz_train = viz_encoded.iloc[:plot_n]
viz_test = viz_encoded.iloc[plot_n:]
viz_y_train = y_sm.loc[viz_train.index]
viz_y_test = y_sm.loc[viz_test.index]
viz_pred = final_sm_model.predict(viz_test)

plt.scatter(viz_train['Dep_Delay'], viz_y_train, color='black', alpha=0.6, u
             label='Train')

```

```

plt.scatter(viz_test['Dep_Delay'], viz_y_test, color='red', alpha=0.6, u
    ↪label='Test')
plt.scatter(viz_test['Dep_Delay'], viz_pred, color='blue', marker='x', s=60, u
    ↪label='Predicted')
plt.title('Regression Analysis: Train vs Test vs Predicted')
plt.xlabel('Departure Delay'); plt.ylabel('Arrival Delay')
plt.legend(); plt.grid(True, alpha=0.3)
plt.show()

```

--- Phase II: Statistical Analysis (Syllabus Requirements) ---

Running Stepwise Regression on 78 features...

OLS Regression Results

```

=====
Dep. Variable:           Arr_Delay      R-squared:                 0.951
Model:                  OLS            Adj. R-squared:             0.951
Method:                 Least Squares  F-statistic:                2.710e+04
Date:                   Fri, 05 Dec 2025  Prob (F-statistic):        0.00
Time:                     18:01:54       Log-Likelihood:          -4.0121e+05
No. Observations:      100000         AIC:                      8.026e+05
Df Residuals:           99928         BIC:                      8.032e+05
Df Model:                   71
Covariance Type:        nonrobust
=====
```

		coef	std err	t	P> t
[0.025	0.975]				
const		-5.8917	0.586	-10.061	0.000
-7.039	-4.744				
Dep_Delay		1.0038	0.001	1367.092	0.000
1.002	1.005				
Flight_Duration		0.0223	0.001	34.017	0.000
0.021	0.024				
Aircraft_age		0.0725	0.006	12.252	0.000
0.061	0.084				
tavg		0.0235	0.005	4.368	0.000
0.013	0.034				
prcp		0.0806	0.006	14.588	0.000
0.070	0.091				
wspd		0.0279	0.008	3.426	0.001
0.012	0.044				
DepTime_label_Evening		-0.8622	0.114	-7.548	0.000
-1.086	-0.638				
DepTime_label_Morning		-0.8458	0.099	-8.528	0.000
-1.040	-0.651				
DepTime_label_Night		-2.2946	0.254	-9.039	0.000

-2.792	-1.797				
Airline_Allegiant Air		4.2173	0.339	12.456	0.000
3.554	4.881				
Airline_American Eagle Airlines Inc.		2.9108	0.258	11.262	0.000
2.404	3.417				
Airline_Delta Air Lines Inc		-2.5763	0.153	-16.816	0.000
-2.877	-2.276				
Airline_Frontier Airlines Inc.		2.0101	0.287	7.000	0.000
1.447	2.573				
Airline_Hawaiian Airlines Inc.		6.7463	0.519	13.000	0.000
5.729	7.763				
Airline_PSA Airlines		2.2957	0.279	8.218	0.000
1.748	2.843				
Airline_Republic Airways		1.4371	0.238	6.033	0.000
0.970	1.904				
Airline_Skywest Airlines Inc.		2.8352	0.177	15.998	0.000
2.488	3.183				
Airline_Southwest Airlines Co.		0.5252	0.138	3.797	0.000
0.254	0.796				
Airline_Spirit Air Lines		0.9996	0.239	4.185	0.000
0.531	1.468				
Airline_United Air Lines Inc.		-1.0995	0.167	-6.601	0.000
-1.426	-0.773				
Origin_State_AL		-3.6798	0.851	-4.325	0.000
-5.347	-2.012				
Origin_State_AR		-4.0970	0.888	-4.614	0.000
-5.837	-2.357				
Origin_State_AS		-16.9533	7.764	-2.184	0.029
-32.170	-1.736				
Origin_State_AZ		-4.9882	0.619	-8.063	0.000
-6.201	-3.776				
Origin_State_CA		-5.0918	0.579	-8.789	0.000
-6.227	-3.956				
Origin_State_CO		-4.2079	0.598	-7.037	0.000
-5.380	-3.036				
Origin_State_CT		-3.9669	0.935	-4.242	0.000
-5.800	-2.134				
Origin_State_FL		-4.3085	0.586	-7.349	0.000
-5.458	-3.159				
Origin_State_GA		-2.6501	0.597	-4.439	0.000
-3.820	-1.480				
Origin_State_HI		-6.1300	0.701	-8.741	0.000
-7.505	-4.755				
Origin_State_IA		-2.8536	0.925	-3.085	0.002
-4.667	-1.040				
Origin_State_ID		-3.3470	0.872	-3.840	0.000
-5.055	-1.639				
Origin_State_IL		-4.6765	0.593	-7.881	0.000

-5.839	-3.513			
Origin_State_IN		-5.6505	0.735	-7.685
-7.092	-4.209			0.000
Origin_State_KS		-3.2238	1.114	-2.895
-5.407	-1.041			0.004
Origin_State_KY		-4.0871	0.699	-5.847
-5.457	-2.717			0.000
Origin_State_LA		-5.1235	0.714	-7.176
-6.523	-3.724			0.000
Origin_State_MA		-6.4139	0.632	-10.149
-7.653	-5.175			0.000
Origin_State_MD		-5.3955	0.670	-8.052
-6.709	-4.082			0.000
Origin_State_ME		-3.3625	1.012	-3.322
-5.346	-1.379			0.001
Origin_State_MI		-5.8031	0.626	-9.271
-7.030	-4.576			0.000
Origin_State_MN		-5.7251	0.646	-8.857
-6.992	-4.458			0.000
Origin_State_MO		-3.6043	0.647	-5.571
-4.872	-2.336			0.000
Origin_State_MS		-3.7463	1.148	-3.264
-5.996	-1.497			0.001
Origin_State_MT		-2.9490	0.894	-3.300
-4.700	-1.198			0.001
Origin_State_NC		-2.8982	0.599	-4.837
-4.072	-1.724			0.000
Origin_State_ND		-6.0061	1.063	-5.648
-8.090	-3.922			0.000
Origin_State_NE		-4.6573	0.911	-5.113
-6.443	-2.872			0.000
Origin_State_NH		-6.0985	1.436	-4.246
-8.914	-3.283			0.000
Origin_State_NJ		-7.4730	0.639	-11.691
-8.726	-6.220			0.000
Origin_State_NM		-2.3103	0.892	-2.589
-4.059	-0.561			0.010
Origin_State_NV		-3.5132	0.618	-5.688
-4.724	-2.303			0.000
Origin_State_NY		-8.2018	0.588	-13.941
-9.355	-7.049			0.000
Origin_State_OH		-4.3317	0.665	-6.510
-5.636	-3.028			0.000
Origin_State_OK		-3.6591	0.800	-4.576
-5.227	-2.092			0.000
Origin_State_OR		-2.6565	0.689	-3.855
-4.007	-1.306			0.000
Origin_State_PA		-5.0021	0.634	-7.890
				0.000

-6.245	-3.760				
Origin_State_PR		-5.3755	0.814	-6.606	0.000
-6.970	-3.781				
Origin_State_RI		-4.6078	1.077	-4.279	0.000
-6.718	-2.497				
Origin_State_SC		-3.7233	0.726	-5.126	0.000
-5.147	-2.300				
Origin_State_SD		-4.4335	1.209	-3.668	0.000
-6.803	-2.064				
Origin_State_TN		-3.8636	0.635	-6.088	0.000
-5.107	-2.620				
Origin_State_TX		-4.1776	0.582	-7.175	0.000
-5.319	-3.036				
Origin_State_UT		-3.3802	0.653	-5.175	0.000
-4.661	-2.100				
Origin_State_VA		-5.9659	0.609	-9.801	0.000
-7.159	-4.773				
Origin_State_VI		-3.7989	1.545	-2.459	0.014
-6.827	-0.771				
Origin_State_VT		-5.2156	1.383	-3.770	0.000
-7.927	-2.504				
Origin_State_WA		-3.0778	0.616	-5.000	0.000
-4.284	-1.871				
Origin_State_WI		-4.9859	0.757	-6.590	0.000
-6.469	-3.503				
Origin_State_WV		-5.4327	1.793	-3.030	0.002
-8.947	-1.918				
Origin_State_WY		-3.0673	1.250	-2.453	0.014
-5.518	-0.617				
<hr/>					
Omnibus:	34575.277	Durbin-Watson:		2.000	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		304305.252	
Skew:	1.416	Prob(JB):		0.00	
Kurtosis:	11.063	Cond. No.		2.94e+04	
<hr/>					

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.94e+04. This might indicate that there are strong multicollinearity or other numerical problems.

	Metric	Value
0	Adjusted R-Squared	0.950599
1	AIC	802554.605750
2	BIC	803239.536383
3	MSE	178.921874

Regression Analysis: Train vs Test vs Predicted

